

Objective-Driven Sequential Bayesian Experimental Design for Power-Grid Dynamic Learning and Control

[TODO: Authors and affiliations]

Abstract—This paper studies sequential Bayesian experimental design for uncertain power-grid dynamics. At each stage, a policy selects a probing bus and duration, observes a noisy frequency response, and updates uncertainty over machine-specific parameters. We distinguish parameter-oriented expected information gain (EIG) from objective-driven mean objective cost of uncertainty (MOCU) for a downstream safety-aware control decision. Random, fixed, myopic, deep adaptive design (DAD), and reinforcement-learning sequential Bayesian optimal experimental design (RL-sBOED) are evaluated under common simulation banks and held-out protocols. A susceptible–infected–recovered (SIR) model provides an implementation benchmark; IEEE 9-bus and IEEE 14-bus systems assess the power-grid application. [TODO: Add final quantitative findings.]

Index Terms—Bayesian experimental design, power-system dynamics, adaptive design, mean objective cost of uncertainty, information gain.

I. INTRODUCTION

UNCERTAINTY in power-system inertia and fast frequency response complicates frequency-security assessment and control selection. Controlled probing can reduce uncertainty, but an experiment must determine where to probe, how long to probe, and whether later probes should depend on earlier observations. [TODO: Add motivation and citations.]

EIG values parameter learning, but parameter information does not necessarily improve a final control decision. We therefore also study MOCU, which values an experiment by its reduction of expected decision penalty.

The contributions are:

- sequential bus–duration BOED for uncertain power-grid dynamics;
- a safety-aware MOCU objective linking probing to terminal control;
- a matched comparison of random, fixed, myopic, DAD, and RL-sBOED; and
- SIR, IEEE9, and IEEE14 studies of objective choice, adaptation, horizon, and scalability.

II. RELATED WORK

Summarize BOED and EIG, adaptive and non-myopic design, DAD, RL-sBOED, MOCU, and power-grid dynamic identification. End with the gap: objective-driven sequential design over power-grid probing buses and durations. [TODO: Add verified references.]

[TODO: Funding and corresponding-author information.]

III. PROBLEM FORMULATION

A. Uncertain Power-Grid Dynamics

For retained dynamic machine i ,

$$\dot{\delta}_i(t) = \omega_i(t), \quad (1)$$

$$M_i \dot{\omega}_i(t) = P_{m,i} - \sum_j B_{ij} \sin(\delta_i - \delta_j) - \left(\frac{K_i}{2\pi} + D_i \right) \omega_i + p_i(t; d). \quad (2)$$

The uncertain vector is

$$\theta = (M_1, \dots, M_G, K_1, \dots, K_G). \quad (3)$$

Only dynamic machines have M_i and K_i . IEEE14 is Kron-reduced onto machines at physical buses 1, 2, 3, 6, and 8. [TODO: State and justify priors.]

B. Sequential Probing Design

At stage t ,

$$d_t = (b_t, \tau_t) \in \mathcal{D}, \quad (4)$$

where b_t is a physical probing bus and τ_t is one of six durations. IEEE9 has $9 \times 6 = 54$ actions and IEEE14 has $14 \times 6 = 84$ actions. The policy selects $d_t = \pi_t(h_{t-1})$ from the previous history.

C. Observation and Belief Update

$$\mathbf{y}_t = g(\theta, d_t) + \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}). \quad (5)$$

Power-grid probes retain five frequency-response observations. Bayesian updating gives

$$p(\theta \mid h_t) \propto p(\mathbf{y}_t \mid \theta, d_t) p(\theta \mid h_{t-1}). \quad (6)$$

D. EIG and MOCU Objectives

$$\text{EIG}(d \mid h) = \mathbb{E}_{\mathbf{y} \mid d, h} [D_{\text{KL}}(p(\theta \mid h, d, \mathbf{y}) \parallel p(\theta \mid h))]. \quad (7)$$

Let $u \in \mathcal{U}$ be a terminal control and $C(u, \theta)$ its safety-aware cost. Define

$$u^*(\theta) = \arg \min_u C(u, \theta), \quad (8)$$

$$u^*(h_T) = \arg \min_u \mathbb{E}_{\theta \mid h_T} [C(u, \theta)], \quad (9)$$

$$\text{OCU}(\theta, h_T) = C(u^*(h_T), \theta) - C(u^*(\theta), \theta), \quad (10)$$

$$\text{MOCU}(h_T) = \mathbb{E}_{\theta \mid h_T} [\text{OCU}(\theta, h_T)]. \quad (11)$$

TABLE I
COMPARED METHODS

Method	Adaptive	Non-myopic	Learned
Random	No	No	No
Fixed	No	Open-loop	No
Myopic	Yes	No	No
DAD	Yes	Yes	Yes
RL-sBOED	Yes	Yes	Yes

The objective-driven policy minimizes $\mathbb{E}_\pi[\text{MOCU}(h_T)]$; lower is better. [TODO: State the implemented safety limits and penalties.]

IV. COMPARED DESIGN METHODS

Random ignores observations. Fixed selects an open-loop sequence. Myopic optimizes immediate expected utility after each observation. DAD learns an amortized history-conditioned policy offline. RL-sBOED learns a sequential policy with stepwise rewards. [TODO: Give reproducible implementation details.]

V. EXPERIMENTAL SETUP

A. SIR EIG Benchmark

SIR verifies the EIG implementation on a simpler nonlinear model; it is not the application contribution. [TODO: State its parameters, actions, and noise.]

B. IEEE9 Experiments

IEEE9 has nine physical probe buses, three dynamic machines, and six uncertain M_i/K_i coordinates. Six durations give 54 actions. Evaluate EIG and MOCU with $N_{\text{obs}} = 5$, $\sigma = 0.005$, $T = 3, \dots, 7$, and seeds 101, 202, and 303. [TODO: Insert final durations and priors.]

C. IEEE14 MOCU Extension

IEEE14 has 14 physical probe buses, five dynamic machines, and ten uncertain M_i/K_i coordinates. Six durations give 84 actions. The MOCU workflow, methods, noise, horizons, and seeds match IEEE9; only IEEE-system physics and machine parameters differ. [TODO: Insert the validated durations.]

D. Duration Selection and Evaluation

Select six durations before policy training using a method-independent search over 281 values from 0.20 to 3.00 s. Screen adaptive value, non-myopic value, and observation-dependent branching; do not tune durations for a reported method. All methods share physical banks, held-out systems, and noise realizations. Report three-seed mean and sample standard deviation, plus paired held-out differences. [TODO: Add bank sizes and evaluation sample counts.]

VI. RESULTS

A. SIR Implementation Validation

[TODO: Report EIG and runtime for every method.]

B. IEEE9 EIG Results

[TODO: Report mean $\pm$ standard deviation for $T = 3, \dots, 7$.]

C. IEEE9 MOCU Results

[TODO: Report terminal MOCU, safety, paired differences, and controls.]

D. IEEE14 MOCU Results

[TODO: Report the same methods and metrics as IEEE9.]

E. Computational Cost

[TODO: Separate bank generation, offline training, and online selection time.]

VII. DISCUSSION

Discuss when adaptation and non-myopic planning improve the downstream objective, why Fixed can be competitive, why EIG and MOCU can disagree, and how horizon and system size affect performance. Limitations include simulated Kron-reduced networks, prescribed priors, fixed probe amplitude, discrete actions, simplified terminal control, and three training seeds. The IEEE14 machine-specific K_i allocation is a modeling assumption, not a published physical value.

VIII. CONCLUSION

This paper formulates power-grid probing as sequential BOED and connects learning to safety-aware terminal control through MOCU. [TODO: State only final conclusions supported by completed three-seed results.]

ACKNOWLEDGMENT

[TODO: Add funding and computing resources.]

REFERENCES

- [1] [TODO: Add verified references for BOED, DAD, RL-sBOED, MOCU, and power-grid dynamic identification.]

TABLE II
 TERMINAL MOCU ACROSS THREE SEEDS (MEAN $\pm$ STANDARD DEVIATION)

Method	$T = 3$	$T = 4$	$T = 5$	$T = 6$	$T = 7$
Random	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
Fixed	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
Myopic	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
DAD	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]
RL-sBOED	[TODO:]	[TODO:]	[TODO:]	[TODO:]	[TODO:]